

## INSTALLATION OF HADOOP IN UBUNTU 14.04 LTS

### MULTI NODE CONFIGURATION

**Step 1. Follow the single node configuration.**

**Configuration done in the master and slave computers. (Assumption: Single Master and Single Slave.) There can be more than 1 slave.**

**Step 2. Note the IP address of master machine and slave machine**

```
$ifconfig
```

**Step 3. In the /etc/hostname file of master add the name of the name-node system.**

```
$sudo gedit /etc/hostname
```

```
master
```

**In the /etc/hostname file of slave add the name of the data-node system.**

```
slave
```

**Step 4. In the /etc/hosts file add the name-node(ip-address, name) and data-nodes(ip-address, name). Name-node is the master and data-node is the slave.**

```
$sudo gedit /etc/hosts
```

```
192.168.1.2 master
```

```
192.168.1.3 slave
```

**Restart the system for changes to take place.**

**Step 5. Configuration of the hadoop files: core-site.xml, mapred-site.xml, hdfs-site.xml and yarn-site.xml**

```
$cd $HADOOP_PREFIX/etc/hadoop
```

verify the path : /usr/local/hadoop-2.5.1/etc/hadoop

**5.1.Configuration of the core-site.xml file**

set the value of the property fs.default.name

```
$sudo gedit core-site.xml  
  
<property>  
<name>fs.default.name</name>  
<value>hdfs://master:53211</value>  
</property>
```

**5.2.Configuration of the mapred-site.xml**

```
$sudo gedit mapred-site.xml
```

set the value of the property mapred.job.tracker

```
<property>  
<name>mapred.job.tracker</name>  
<value>master:53212</value>  
</property>
```

**5.3.Configuration of the hdfs-site.xml**

```
$sudo gedit hdfs-site.xml
```

set the value of the properties dfs.replication and dfs.data.dir as

```
<property>
<name>dfs.replication</name>
<value>2</value>
</property>
<property>
<name>dfs.data.dir</name>
<value>/usr/local/hadoop/hdfs</value>
</property>
```

set the value of the properties yarn.nodemanager.auxservices,yarn.nodemanager.aux-services.

mapreduce\_shuffle.class, yarn.resourcemanager.resource-tracker.address,  
yarn.resourcemanager.scheduler.address,yarn.resourcemanager.address

#### 5.4.Configuration of theyarn-site.xml

```
$sudo gedit yarn-site.xml

<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.aux-
services.mapreduce_shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<property>
<name>yarn.resourcemanager.resource-tracker.address</name>
<value>master:8025</value>
</property>
```

```
<property>
<name>yarn.resourcemanager.scheduler.address</name>
<value>master:8030</value>
</property>
<property>
<name>yarn.resourcemanager.address</name>
<value>master:8050</value>
</property>
```

### Configuration done in the master computer

**Step 6. Enable ssh so that master can access all slaves including itself.**

```
$ssh-copy-id -i ~/.ssh/id_rsa.pub hdptools@master
$ssh-copy-id -i ~/.ssh/id_rsa.pub hdptools@slave
```

**Step 7. Test ssh connection**

**Check master connection**

```
$ssh master
```

**After successful ssh connection.**

```
$exit
```

**Check Slave connection**

```
$ssh slave
```

```
$exit
```

**Step 8. In \$HADOOP\_PREFIX/etc/hadoop/masters, set the name-node(master) name.**

```
$sudo gedit masters
```

```
master
```

7.3. In \$HADOOP\_PREFIX/etc/hadoop/slaves, set the data-node(slave) names.(master can act also as one of the data-node)

```
$sudo gedit slaves
```

```
master
```

```
slave
```

8.Format the Hadoop File system implemented on top of the local file system using

```
$cd bin
```

Verify the path : /usr/local/hadoop-2.5.1/bin

```
$hadoopnamenode -format
```

```
$cd ..
```

9. In the master, Start Hadoop using

```
$cd sbin
```

Verify the path : /usr/local/hadoop-2.5.1/sbin

```
$./start-all.sh
```

```
$jps
```

10. Accessing Hadoop on Browser

The default port number to access Hadoop is 50070. Use the following url to get Hadoop services on browser.

<http://localhost:50070/>

## 11. Verify All Applications for Cluster

The default port number to access all applications of cluster is 8088. Use the following url to visit this service.

<http://localhost:8088/>

## COMPILATION AND EXECUTION OF MAP REDUCE PROGRAM

1. Write a Map Reduce program and save it in the \$HOME Path.

Eg.:WordCount.java

2. Create a empty directory “WordCount” in \$HOME path. The following command is to create a directory to store the compiled java classes.

```
cd $HOME
```

```
mkdir WordCount
```

3. Download Hadoop-core-1.2.1.jar, which is used to compile and execute the Map Reduce program.

Visit the following link:

<http://mvnrepository.com/artifact/org.apache.hadoop/hadoop-core/1.2.1>

to download the jar. Let us assume the downloaded folder is

[/usr/local/hadoop-2.5.1/](#)

4. Compile the java program by specifying the class path:

```
javac -classpath /usr/local/hadoop-2.5.1/hadoop-core-1.2.1.jar -d WordCount WordCount.java
```

5. Create jar files for the program

```
jar -cvf WordCount.jar -C WordCount/
```

6. Create an input file “fsample.txt” contains some text in \$HOME path.

7. Change directory to Hadoop

```
cd $HADOOP_PREFIX
```

8. The following command is used to create an input directory in HDFS.

```
bin/hdfs dfs -mkdir input_dir
```

9. The following command is used to copy the input file named **fsample.txt** in the input directory of HDFS.

```
bin/hdfs dfs -put $HOME/fsample.txt /input_dir
```

10. The following command is used to verify the files in the input directory.

```
bin/hdfs dfs -ls input_dir/
```

11. The following command is used to run the application by taking the input files from the input directory.

```
bin/hadoop jar $HOME/WordCount.jar hadoop.WordCount  
/input_dir /output_dir
```

12. The following command is used to verify the resultant files in the output folder.

```
bin/hdfs dfs -ls output_dir/
```

13. The following command is used to see the output in **Part-00000** file. This file is generated by HDFS.

```
bin/hdfs dfs -cat output_dir/part-00000
```

14. Copy the output files from the distributed file system to the local file system and examine them:

```
/usr/local/hadoop-2.5.1$ bin/hdfs dfs -get output/*  
/home/huser/Downloads/output/
```

```
/usr/local/hadoop-2.5.1$ cat output/*
```

---